In this course, we will work through most of Speech and Language Processing by Jurafsky and Martin. In principle, we follow the Speech + NLP track in the preface. However, we have added some information from other chapters that introduce concepts we deem important. To compensate, we dropped some less important parts from the track chapters.

The Book

You will need the book. We use the Second Edition. On first sight, it looks impressively thick (and you will be expected to read quite a few pages every week). But it is written quite well, explaining all concepts extensively and giving examples, so that we think the amount of reading is certainly not too much.

In case you don't have the book yet, we will provide PDFs for a draft version of the Second Edition so you can participate. Beware, however, that some details (hopefully only unimportant ones) may not yet be in the definitive form (as found in the book). We already know that numbering of exercises is sometimes different.

If you happen to find a copy of the first edition on Ebay or so, do ***not*** buy (at least for this course). The Second Edition is a significant improvement.

Exams

The exams will be open-book-exams, i.e. you are allowed to bring the book and any other materials you want (except means of communication) and use them at the exam. This fact also influences the way you will be reading the book: you will need to be able to know your way around and to understand what is there. You do not need to know everything by heart. However, you should know by heart the main ideas and terminology, as there will not be enough time to look up everything.

Activities

During the course, we do the following:

- The week before a chapter should be read, I give a short introduction and point out any problematic parts in the chapter.
- Between then and the class, you read and do your best to understand the chapter.
- You also do the prescribed exercises.
- In class, I first sample the preparedness of the students present (which includes inspecting the exercise results); if you are not prepared, you are marked as not present (although you are allowed to stay during the class)
- Then (most of the class time) we go through the chapter(s) and exercises together, resolving any problems you may have understanding the topics.

Important: You need to be marked as present at least 11 out of the 14 weeks. In order to be marked as present, you need to be physically present during the whole class and participate actively. You also need to be prepared, by having read the chapter(s) for that week and having completed the prescribed exercises.

Reading levels

In the Blackboard section for every week, we will list what you are supposed to read and at what level:

- R+U. Read and understand. So read at full capacity, but you are not required to know everything by heart, as mentioned above.
- R. Just read, in order to get a feel for the background and what else exists.

Apart from these listed instructions, the following general rules are in effect:

- All summaries: you should attempt to know this information by heart (so you can talk to others about these topics without looking things up all the time).
- All terms (bolded words in the text): you have to know what these represent by heart.
- All algorithms and formulas: you have understand them and their use, and to know where to find the details about them. For some sections, we will signal that this is not necessary.
- All Bibliographical and Historical Notes: Level R: you have to read this to get a feel for the background and what else exists.

General Schedule

In the following table, we list how we spread out the book over the weeks of the course. Every week, we will list which sections of the chapters to read and at what level.

Week	Chapters	Estimated pages RU+R
1	1, 2, 3 (without 3.10, 3.11)	40+15
2	4, 5 (without 5.9)	45+30
3	6, Intro to Machine Learning	35+20
4	7, 8	50
5	9	50
6	12, 16	40
7	13, 15	50
Exam	Midterm	
8	(Look back at period 1); 3.10, 3.11; 5.9; 14	35+5
9	17, 18	40
10	19, 21	35
11	20, Intro to word2vec/GloVe	55
12	22, 23	60
13	24	50
14	25	50
Exam	Final	

As stated above, do not panic about the number of pages. The book is quite friendly.

Exercises per week should normally be doable within one to two hours. Given the tight schedule, we will avoid exercises in which algorithms need to be implemented.